AI MASTERY GUIDE

From Fundamentals to Real-World Fintech AI Applications

1. AI Fundamentals	**2. AI for Brainstorming**	**3. AI for Research**	**4. AI for Writing**
5. AI for Content Creation	**6. AI for Data Analysis**	**7. AI for App Building**	

Written in simple English. Packed with Fiserv-specific examples. Designed to help you implement AI with confidence.

Edition: February 2026 | ~50 Pages | 7 Chapters

TABLE OF CONTENTS

INTRODUCTION

Why AI Matters at Fiserv

Fiserv is one of the largest fintech companies in the world. It processes billions of financial transactions every year. Think of credit-card swipes, bank transfers, loan approvals, and mobile payments — a huge share of those go through Fiserv's systems. As an AI Implementation Engineer here, your job is to make those systems smarter, faster, and safer by using Artificial Intelligence.

This guide is your starting point. It is written in plain, everyday English. You do not need a PhD in mathematics to understand it. Every chapter covers one AI skill area, explains the concept simply, then shows you exactly how to use it in the Fiserv context.

What is the 'AI Divide'?

You may have heard people say that AI will replace workers. That is not quite right. What is actually happening is this: workers who know how to use AI are becoming far more productive than those who do not. The real gap is not about intelligence — it is about knowing where to start.

This guide gives you exactly that: a clear starting point across seven key skill areas.

How to Use This Guide

- Read one chapter per day. Each chapter takes about 20–30 minutes.
- Try the examples immediately. Open Claude, ChatGPT, or Copilot and paste the prompts shown.
- Adapt every example to a real Fiserv project you are working on.
- Share what you learn with your team. AI skills grow faster in teams.

■ **Quick Tip**	Every chapter ends with a 'Fiserv Fast Start' — a ready-to-use prompt or action you can apply to your work today. Do not skip these!

CHAPTER 1

AI Fundamentals

Understanding What AI Is, How It Thinks, and Where It Fits in Fintech

1.1 What Is Artificial Intelligence, Really?

Artificial Intelligence (AI) is software that can do things we usually need human intelligence for — like understanding language, recognising patterns, making predictions, or generating new content. It learns from huge amounts of data instead of being programmed rule by rule.

Think of it this way: a traditional banking system has strict rules written by developers. If a transaction is above $10,000, flag it. Period. An AI system looks at thousands of signals — time of day, location, spending history, merchant type — and calculates a fraud probability score. It learns from every transaction, improving over time.

Three Types of AI You Will Work With

Machine Learning (ML)	AI that learns patterns from data. Example: Fiserv's fraud detection model learns from millions of past fraudulent transactions.
Natural Language Processing (NLP)	AI that reads, understands, and generates human language. Example: a chatbot that answers customer banking questions.
Generative AI (GenAI)	AI that creates new content — text, code, images. Example: Claude or GPT-4 helping you write an incident report or generate code.

1.2 Key AI Vocabulary (No Jargon!)

Before diving into the chapters, here are the 12 words you will hear every day as an AI Implementation Engineer at Fiserv. Understanding these is like knowing the alphabet — everything else becomes much easier.

Model	The trained AI brain. A model is a file that has learned patterns from data and can now make predictions or generate text.
Training	The process of feeding data to an AI so it can learn. Like teaching a new employee by showing them thousands of examples.

Inference	Using a trained model to answer a new question or process new data. When a customer swipes a card, fraud inference happens in milliseconds.
Prompt	The instruction or question you give to a Generative AI. Think of it as talking to a very knowledgeable assistant — what you say determines what you get back.
Token	The basic unit of text that AI processes. Roughly, 1 token ≈ 0.75 words. Models have a maximum number of tokens they can read at once (their 'context window').
Embedding	A way to turn text into numbers so AI can compare meaning. 'Payment failed' and 'transaction declined' have very similar embeddings because they mean the same thing.
Vector Database	A special database that stores embeddings. Used in RAG systems (see below) to find the most relevant documents quickly.
RAG	Retrieval-Augmented Generation. AI that searches your company's documents first, then answers based on what it finds. Great for Fiserv's internal knowledge bases.
Hallucination	When AI confidently states something that is wrong. Always verify AI output before using it in production financial systems.
Fine-tuning	Retraining a general AI model on your specific data so it becomes expert in your domain — e.g., training a model on Fiserv's payment terminology.
API	Application Programming Interface. The way your code talks to an AI service like Azure OpenAI or AWS Bedrock. You send a request, you get a response.
Orchestration	Connecting multiple AI steps together. For example: user asks a question → AI searches documents → AI generates answer → system sends reply. Frameworks like LangChain help with this.

1.3 How Large Language Models (LLMs) Work

Most Generative AI tools you will use (Claude, GPT-4, Gemini) are Large Language Models. Here is a simple explanation of how they work, without the complex maths:

- Step 1 — Read everything: The model was trained on a massive amount of text from the internet, books, and code. This is how it learned grammar, facts, and reasoning.
- Step 2 — Learn patterns: During training, the model learned that certain words and ideas follow certain patterns. After 'The customer's credit card was', the word 'declined' is more likely than 'blue'.
- Step 3 — Predict the next word: When you ask it a question, it generates a response one word (token) at a time, always choosing the most sensible next word based on everything before it.
- Step 4 — Follow instructions: Modern models are also trained to follow instructions helpfully and safely — this is called RLHF (Reinforcement Learning from Human Feedback).

■ **Fiserv Context**	When you use Azure OpenAI at Fiserv, your prompts stay within Fiserv's Azure tenant. The data you send does NOT train the public model. This is a critical security guarantee for financial data compliance (PCI-DSS, SOX).

1.4 The AI Stack at Fiserv

As an AI Implementation Engineer, you will work across multiple layers of an AI system. Here is how the typical Fiserv AI stack looks, from bottom to top:

Layer 5 — Applications	Customer-facing chatbots, internal tools, fraud dashboards, developer copilots
Layer 4 — Orchestration	LangChain, LlamaIndex, AutoGen — connect AI models with tools and data
Layer 3 — AI Models	Azure OpenAI (GPT-4), Anthropic Claude, AWS Bedrock, custom ML models
Layer 2 — Data & Storage	Pinecone vector DB, Azure Cosmos DB, Fiserv data lakes, Kafka streams
Layer 1 — Infrastructure	Azure, AWS, on-prem servers, Kubernetes, Docker, CI/CD pipelines

1.5 AI Safety and Compliance in Fintech

Financial services are heavily regulated. As an AI Implementation Engineer at Fiserv, you must understand these non-negotiables:

- Explainability: Regulators (OCC, CFPB) require that AI decisions affecting customers can be explained. A black-box model that denies a loan is not acceptable without an explanation.
- Bias testing: AI models must be tested for unfair bias across race, gender, and age before deployment in credit or lending systems.

• Data privacy: PII (Personally Identifiable Information) must never be sent to external AI APIs without proper data handling agreements. Use anonymisation or synthetic data for testing.

• Audit trails: Every AI decision in a regulated workflow must be logged. Who asked? What was the input? What did the model output? When?

• Human oversight: For high-stakes decisions (loan approvals, fraud disputes), AI is an assistant to a human decision-maker, not a replacement.

■ Fiserv Fast Start — Chapter 1	Open your AI tool and try this prompt: 'I am an AI Implementation Engineer at a fintech company. Explain how a fraud detection ML model works, using a simple analogy. Then list 5 ways it is different from a traditional rules-based fraud system.' Study the response and notice how clearly AI can explain complex technical concepts.

CHAPTER 2

AI for Brainstorming

Generate Ideas at Scale — Sprint Planning, Architecture, and Innovation Sessions

2.1 Why Brainstorming Is Different with AI

Traditional brainstorming relies on a room of people. You get maybe 20–30 ideas in an hour. AI brainstorming gives you 50–100 ideas in 5 minutes, from multiple angles, without groupthink. Your job then is to evaluate and refine — which is a much more productive use of your expertise.

For an AI Implementation Engineer, brainstorming tasks come up constantly: What features should our new payment AI assistant have? What edge cases should we test for? How should we architect this RAG system? AI is your always-available thought partner.

2.2 The Anatomy of a Great Brainstorming Prompt

Most people get mediocre AI brainstorming results because their prompts are vague. Here is a simple formula that works every time:

Role	Tell the AI who it is. 'You are a senior fintech architect with 15 years of experience.'
Context	Give background. 'We are building a real-time payment fraud detection system for a Tier-1 bank.'
Task	Be specific. 'Brainstorm 10 potential failure modes we should test for before going live.'
Format	Specify output. 'List each failure mode with: name, description, severity (High/Medium/Low), and how to test it.'
Constraints	Add limits. 'Focus on failures that could occur within the first 30 days of production.'

2.3 10 Brainstorming Use Cases for Fiserv Engineers

Use Case 1: Architecture Design

Before designing a new AI system, use AI to brainstorm architectural options, trade-offs, and potential problems before you write a single line of code.

Example Prompt:

```
"We need to build a document intelligence system that can read merchant
contracts (PDFs), extract key terms (payment rates, chargeback limits,
termination clauses), and flag unusual clauses for legal review. Brainstorm 5
different architectural approaches, with pros and cons of each. Consider:
accuracy, cost, speed, and compliance."
```

Use Case 2: Test Case Generation

Writing comprehensive test cases is tedious. AI can generate hundreds of edge cases you might never have thought of.

Example Prompt:

```
"Generate 20 test cases for a payment fraud detection API. Include: normal
transactions, edge cases (exactly at threshold amounts), unusual merchant
categories, cross-border transactions, and known fraud patterns. Format as a
table with: Test ID, Scenario, Input Data, Expected Output."
```

Use Case 3: Incident Post-Mortem Analysis

After a production incident, AI can help you think through root causes and prevention strategies systematically, ensuring you do not miss anything in the heat of the moment.

Use Case 4: Sprint Planning Ideas

AI can suggest user stories, acceptance criteria, and technical tasks for a sprint based on a high-level feature description — cutting planning time significantly.

Example Prompt:

```
"We are planning a 2-week sprint to improve our AI fraud model's explainability.
Generate 8 user stories with acceptance criteria. The stakeholders are:
compliance team, customer service agents, and backend engineers. Use standard
user story format."
```

Use Case 5: API Design Brainstorming

Before designing a new AI service endpoint, brainstorm the ideal API contract, including request/response schemas, error codes, and rate limiting strategies.

Use Case 6: Risk Identification

For any new AI deployment, use AI to brainstorm risks across technical, compliance, operational, and reputational dimensions. It is remarkably thorough.

Example Prompt:

```
"We are about to deploy an AI chatbot that can answer customer questions about
their bank account. Brainstorm 15 risks we need to mitigate before launch.
Categorize by: Security, Compliance, Accuracy, Customer Experience, and
```

```
Operational."
```

Use Case 7: Naming and Messaging

Need to name a new AI product feature or internal tool? AI can generate dozens of options in seconds, each with a brief rationale.

Use Case 8: Technology Evaluation Criteria

When evaluating two different AI tools (e.g., Pinecone vs Weaviate for vector storage), AI can help you brainstorm a comprehensive evaluation matrix tailored to your specific needs.

Use Case 9: On-call Runbook Ideas

AI can help you brainstorm what to include in runbooks for AI system incidents — what metrics to check, what actions to take, who to escalate to.

Use Case 10: Training Content Ideas

Need to train the rest of your team on a new AI capability? AI can brainstorm a curriculum, hands-on exercises, quizzes, and real-world scenarios in minutes.

2.4 Advanced Brainstorming Techniques

Technique 1: Perspective Shifting

Ask AI to brainstorm from the perspective of different stakeholders. The same feature looks very different to a compliance officer vs a developer vs a customer.

```
"Brainstorm concerns about our new AI loan-decisioning system from the
perspective of: a) a bank regulator, b) a loan applicant who was rejected, c)
the bank's Chief Risk Officer, d) the DevOps engineer who has to maintain it."
```

Technique 2: Reverse Brainstorming

Instead of asking 'How do we make this better?', ask 'How could we make this catastrophically bad?' Then reverse each answer into a positive design principle. This is surprisingly effective for finding blind spots.

```
"List 10 ways we could make our fraud detection AI completely useless or
harmful. For each one, state the corresponding design principle that prevents
it."
```

Technique 3: Quantity then Quality

First ask for 20 ideas quickly, without detail. Then pick the 3 most interesting and ask AI to deep-dive on just those. This prevents you from getting stuck on the first idea you see.

■ Fiserv Fast Start — Chapter 2	Try this right now: 'List 15 ways an AI Implementation Engineer at a payment processing company could use Generative AI to reduce manual work in their daily job. Be specific to fintech workflows.' Pick your top 3 ideas and implement at least one this week.

CHAPTER 3

AI for Research

Find Answers Faster — Technology Evaluation, Competitive Analysis, and Learning New Topics

3.1 The Research Problem in Fast-Moving AI

The AI field moves incredibly fast. New models, frameworks, and tools appear every month. As an AI Implementation Engineer, you need to stay current, evaluate new technologies quickly, and make evidence-based recommendations. AI dramatically accelerates all of this.

Traditional research: You spend 3 days reading documentation, blog posts, and papers to understand a new tool. With AI: you spend 30 minutes having a conversation that gives you a deep mental model, then use the remaining time for hands-on testing.

3.2 How to Use AI as a Research Partner

Step 1: Orient Yourself

When you encounter a completely new topic, start with a high-level overview prompt to build your mental model before going deep.

```
"I am an AI engineer who needs to understand GraphRAG (Graph-based Retrieval
Augmented Generation) quickly. Explain: 1) What it is in simple terms, 2) How it
differs from standard RAG, 3) When you would use it vs standard RAG, 4) The main
technical components, 5) The main limitations. Keep each section to 3-4
sentences."
```

Step 2: Go Deep on Specifics

Once you have the overview, drill into the specific aspects most relevant to your work.

```
"Now explain specifically how GraphRAG handles entity disambiguation in
financial documents where the same company name might refer to different legal
entities. What are the specific challenges and how does the graph structure
help?"
```

Step 3: Connect to Your Context

Finally, make the research directly applicable to your work.

```
"Given what you have explained about GraphRAG, design a high-level architecture
for using it to search Fiserv's merchant agreement documents, where merchants
often have complex corporate structures. What would the graph nodes and edges
represent?"
```

3.3 Research Use Cases for Fiserv Engineers

Technology Evaluation

When your team needs to choose between two technology options, AI can rapidly synthesise a comparison based on your specific requirements.

```
"Compare Pinecone, Weaviate, and pgvector for a Fiserv use case: we need to
store embeddings of 50 million financial documents, support real-time updates,
and query with low latency (under 100ms) at 1,000 queries per second. We are
already on Azure. Give a recommendation with justification."
```

Regulatory Research

Understanding how regulations apply to AI systems is complex. AI can help you parse regulatory requirements and translate them into technical implications.

```
"Summarise the key requirements of the EU AI Act that apply to an AI-powered
credit scoring system used by a US-based bank operating in Europe. Focus on:
transparency requirements, human oversight, and technical documentation
obligations."
```

Competitor Analysis

Understanding how competitors are using AI can inform your own strategy and help you identify gaps and opportunities.

Learning a New Framework

When your team adopts a new framework (e.g., LangGraph for agentic workflows), AI can create a customised learning path tailored to your existing knowledge.

```
"I know Python, FastAPI, and basic LangChain. Create a 5-day learning plan for
me to become productive with LangGraph for building multi-agent AI systems. For
each day, list: topic, key concepts, a hands-on exercise, and resources. Focus
on patterns relevant to financial services workflows."
```

3.4 Research Best Practices and Pitfalls

Always Verify AI Research

AI can make mistakes, especially on very recent developments or highly specific technical details. Treat AI research as a starting point, not a final answer. Key verification steps:

- For technical facts: check the official documentation of the technology.
- For regulatory information: always consult Fiserv's legal or compliance team for final guidance.
- For model capabilities: test with actual code — AI descriptions of performance are approximate.
- For recent events: AI's knowledge has a cutoff date; use web search for anything from the last 6 months.

The Research Conversation Pattern

The best AI research is a conversation, not a single query. A typical pattern:

- Start broad → get the overview
- Ask follow-up questions on the confusing parts
- Challenge the AI — 'What are the weaknesses of this approach?'
- Ask for alternatives — 'What would a sceptic say about this?'
- Synthesise — 'Summarise the 3 most important things I should remember from our conversation'

■ Fiserv Fast Start — Chapter 3	Research challenge: Ask AI to compare the top 3 vector databases for a payment data use case, specifying that you need GDPR compliance, under 50ms query time, and Azure deployment. Then take that output and verify one key claim against the official documentation. This builds the habit of AI-assisted-but-verified research.

CHAPTER 4

AI for Writing

Produce Clear, Professional Documents — Technical Specs, Reports, and Communications

4.1 Writing Is a Core Engineering Skill

Great engineers are great communicators. As an AI Implementation Engineer, you will write: technical design documents, architecture decision records, incident reports, business cases, API documentation, stakeholder updates, and much more. Poor writing creates confusion, slows decisions, and undermines your credibility. AI makes you a significantly better writer.

4.2 The Writing Workflow with AI

The most effective way to use AI for writing is not to have it write everything from scratch. It is a collaborative process:

1. Draft your thoughts	Write a rough, messy first version. Even just bullet points. Your knowledge, your ideas.
2. Hand to AI for structure	'Turn these bullet points into a structured technical document with an executive summary.'
3. Refine together	Ask AI to improve specific sections, adjust tone, add missing details.
4. You review and own it	Always read the final version carefully. Add your specific knowledge. Remove anything incorrect.

4.3 Essential Writing Templates for Fiserv Engineers

Template 1: Technical Design Document (TDD)

Use this prompt to scaffold a TDD for any new AI feature:

```
"Write a Technical Design Document for the following AI system: [describe your
system]. Include these sections: 1) Executive Summary (3 sentences), 2) Problem
Statement, 3) Proposed Solution Architecture (with component descriptions), 4)
Data Flow, 5) API Contracts (request/response examples), 6) Security and
Compliance Considerations, 7) Performance Requirements, 8) Testing Strategy, 9)
Rollout Plan, 10) Open Questions."
```

Template 2: Incident Post-Mortem Report

After any production incident with an AI system, use this to write a thorough post-mortem:

```
"Write an incident post-mortem report based on these facts: [paste your
timeline/notes]. Use this structure: Incident Summary, Impact (customers
affected, duration, financial impact), Root Cause Analysis (5 Whys format), What
Went Well, What Went Wrong, Action Items with owners and deadlines. Use
professional, blameless language."
```

Template 3: Business Case for AI Investment

When you need to justify an AI project to leadership, use this structure:

```
"Write a business case for the following AI initiative: [describe initiative].
Include: Problem and Opportunity (with data), Proposed Solution, Expected
Benefits (quantified where possible), Implementation Costs (rough order of
magnitude), Timeline, Risks and Mitigations, Recommended Decision. Tone should
be executive-friendly — clear, concise, data-driven."
```

Template 4: Architecture Decision Record (ADR)

ADRs document why you made a technical decision — invaluable for future engineers and auditors.

```
"Write an Architecture Decision Record for this decision: [describe decision].
Format: Title, Date, Status (Proposed/Accepted/Deprecated), Context (why we
needed to decide), Decision Options Considered (with pros/cons table), Decision
Made and Rationale, Consequences (positive and negative)."
```

Template 5: Stakeholder Email/Update

For non-technical stakeholders, translate your technical work into business impact:

```
"Write a stakeholder update email about this AI project progress: [your bullet
points]. Audience: VP of Product and business stakeholders (non-technical).
Tone: Professional, confident, concise. Structure: 1-line subject, 3-sentence
summary, What we accomplished this week, Key risks/blockers, What we're doing
next week, What we need from them (if anything)."
```

4.4 Making AI Writing Sound Like You

AI writing can sometimes sound generic. Here is how to make it sound like you:

- Feed AI your previous writing: 'Here is an email I wrote. Match this writing style: [paste email]'
- Be specific about tone: 'Write this formally but avoid corporate jargon. Direct and confident.'
- Add your examples: After AI drafts, add your own specific examples and data points.
- Read it aloud: If it sounds unnatural when you say it, fix it. You own the document.
- Personality prompt: 'I am pragmatic, detail-oriented, and direct. Write in that voice.'

4.5 API Documentation That Developers Actually Love

When you build AI services at Fiserv, other developers will consume your APIs. Good documentation is essential. Here is a prompt that generates great API docs:

```
"Write API documentation for this endpoint: [describe your endpoint, parameters,
response]. Include: Overview paragraph, Authentication requirements, Request
parameters table (name/type/required/description), Example request (curl
format), Success response (JSON), Error responses table, Rate limiting
information, and a code example in Python."
```

■ Fiserv Fast Start — Chapter 4	Take a document you wrote recently (an email, a Jira ticket, a Confluence page). Paste it into AI and ask: 'Rewrite this to be 30% shorter while keeping all key information. Then suggest 3 ways to make the structure clearer.' Compare the result to your original. This exercise will teach you more about writing than any style guide.

CHAPTER 5

AI for Content Creation

Build Training Materials, Demos, Presentations, and Knowledge Assets at Scale

5.1 Content Creation in the AI Era

As an AI Implementation Engineer, you create more content than you might realise: onboarding documentation for new team members, presentations for leadership, training materials for users of your AI systems, demo scripts, and knowledge base articles. AI can multiply your output dramatically — but the quality still depends on your expertise and judgment.

5.2 Training Materials for AI System Users

When you deploy a new AI tool at Fiserv — say, an AI assistant for customer service agents — you need to train the users. Here is how AI helps you create that training:

Creating a Training Curriculum

```
"Design a 3-hour training programme for Fiserv customer service agents who will
use a new AI assistant that can look up account information and answer policy
questions. Include: learning objectives, agenda with timing, hands-on
exercises, common mistakes to avoid, and an assessment quiz with 10 questions."
```

Writing User Guides

```
"Write a user guide for customer service agents explaining how to use our new AI
assistant. Audience: agents with no technical background. Tone: friendly,
step-by-step. Include: what the tool can and cannot do, how to phrase questions
to get the best answers, when to use the tool vs escalate to a human, and what
to do if the AI gives wrong information."
```

Creating FAQ Documents

```
"Generate a 20-question FAQ for employees asking about Fiserv's new AI fraud
detection system. Cover: how it works, why it was implemented, what happens when
it flags a transaction, how accuracy is measured, and privacy/data handling.
Answers should be 2-4 sentences each."
```

5.3 Presentation Creation

Executive Presentations

Executives need the big picture, the business impact, and the risks — not technical details. AI helps you translate your technical work into executive language.

```
"Create a presentation outline for a 15-minute executive briefing on our AI
fraud detection project. Audience: C-suite and VPs. Include: slide titles, 3
bullet points per slide, and a 'so what' statement for each slide. Start with
business impact, end with ask/recommendation."
```

Technical Deep-Dives

```
"Create a 30-minute technical presentation for our engineering team on RAG
(Retrieval-Augmented Generation) architecture. Include: concept explanation,
architecture diagram description, demo ideas, code walkthrough outline, and
common pitfalls. Assume the audience knows Python and REST APIs but is new to
vector databases."
```

5.4 Demo and Prototype Content

When you are building a proof of concept or demo of a new AI system, you need realistic sample data that looks authentic but is not real customer data. AI excels at generating this:

```
"Generate 20 realistic synthetic customer service chat conversations between a
bank customer and an AI assistant. Scenarios: account balance inquiry, failed
transaction dispute, lost card report, loan status question, and general banking
question. Each conversation should be 4-6 exchanges. Make the language realistic
— customers are not always polite or clear."
```

5.5 Knowledge Base Articles

Your team's Confluence, SharePoint, or internal wiki is only valuable if it has good content. AI can help you rapidly create and maintain knowledge base articles:

- Runbooks: 'Write a runbook for when our fraud detection model accuracy drops below 90%. Include: how to detect it, initial diagnosis steps, escalation path, and rollback procedure.'
- How-to guides: 'Write a how-to guide for a junior engineer joining our team to set up their local development environment for our AI services.'
- Glossary articles: 'Write a glossary of 20 AI terms for Fiserv business stakeholders who are not technical, with examples from the payments industry.'
- Post-mortems as learning content: 'Convert this incident post-mortem into a 5-minute learning article that teaches the key lessons to future engineers.'

5.6 Content Quality Control

AI-generated content needs review. Use AI itself to quality-check its output:

```
"Review the training document I just created (paste content). Check for:
accuracy of technical claims, clarity for a non-technical audience, missing
important topics, inconsistencies, and overly complex language. Suggest
specific improvements."
```

■ Fiserv Fast Start — Chapter 5	Create a one-page 'AI Tools Cheat Sheet' for your team using AI. Prompt: 'Create a one-page cheat sheet for software engineers at a fintech company showing how to use AI tools in their daily work. Cover: code review, debugging, documentation, testing, and learning. Format as a clean reference card with examples.' Share it with your team and collect feedback.

CHAPTER 6

AI for Data Analysis

Turn Raw Data into Decisions — Metrics, Monitoring, and Model Evaluation

6.1 Data Analysis in Fintech AI

At Fiserv, data is the lifeblood of AI. Every AI model you build is only as good as the data it learns from and the metrics you use to evaluate it. As an AI Implementation Engineer, you will constantly analyse: model performance metrics, transaction data patterns, system health data, A/B test results, and business impact metrics. AI dramatically speeds up this analytical work.

6.2 Using AI to Write Data Analysis Code

Exploratory Data Analysis (EDA)

Instead of writing EDA code from scratch, describe your dataset to AI and get ready-to-run code:

```
"Write Python code to perform exploratory data analysis on a CSV file containing
payment transaction data with these columns: transaction_id, amount,
merchant_category, timestamp, customer_id, is_fraud (boolean), country_code.
Include: 1) Basic statistics summary, 2) Fraud rate by merchant category (bar
chart), 3) Transaction amount distribution for fraud vs non-fraud (overlapping
histogram), 4) Fraud rate by hour of day, 5) Missing value analysis. Use pandas
and matplotlib."
```

SQL Query Generation

Describe what you want to know from your database and AI writes the SQL:

```
"Write a SQL query to find merchants where the fraud rate increased by more than
20% in the last 30 days compared to the prior 30 days. The transactions table
has: transaction_id, merchant_id, amount, timestamp, is_fraud. Return:
merchant_id, prior_period_fraud_rate, recent_fraud_rate, change_percentage.
Order by change_percentage descending."
```

Data Pipeline Code

```
"Write a Python function that reads transaction data from an Azure Blob Storage
CSV file, computes daily fraud metrics (transaction count, fraud count, fraud
rate, average fraud amount), and saves results to Azure Cosmos DB. Include error
handling and logging. Use Azure SDK for Python."
```

6.3 AI Model Evaluation

Evaluating an AI model's performance is one of the most important skills for an AI Implementation Engineer. Here is a framework and how AI helps you apply it:

Accuracy	What % of predictions are correct? For fraud: (True Positives + True Negatives) / Total. Use AI to: generate the evaluation code, interpret the results.
Precision	Of the transactions we flagged as fraud, what % actually were fraud? Low precision = too many false alarms (customer frustration).
Recall	Of all actual fraud transactions, what % did we catch? Low recall = missed fraud (financial loss).
F1 Score	Harmonic mean of precision and recall. Use when you need to balance both. AI can explain the trade-offs for your specific use case.
AUC-ROC	How well does the model distinguish fraud from non-fraud across all thresholds? AI can generate the ROC curve plot and explain what the AUC score means.
Business Metrics	Cost of missed fraud vs cost of false alarms. AI can help you build a cost-benefit model to find the optimal decision threshold.

Automated Model Evaluation Prompt

```
"Write Python code to evaluate a binary fraud classification model. Input:
predictions CSV with columns 'actual' and 'predicted_probability'. Compute and
display: accuracy, precision, recall, F1 score, AUC-ROC with ROC curve plot,
confusion matrix heatmap, and precision-recall curve. Also compute: the optimal
threshold to maximise F1 score, and the business impact if each false negative
costs $500 and each false positive costs $10."
```

6.4 Interpreting Analysis Results

Once you have your analysis results, AI can help you interpret them and decide what to do:

```
"I ran a fraud detection model evaluation and got these results: Precision:
0.87, Recall: 0.71, AUC: 0.94, F1: 0.78. The model is deployed in a consumer
debit card product with 500,000 transactions per day. Interpret these metrics in
plain English. Is this good or bad? What does the recall score mean for
customers? What actions should I consider to improve performance?"
```

6.5 Monitoring AI Systems in Production

Deploying a model is just the beginning. AI models can degrade over time as real-world patterns change — this is called 'model drift'. Here is how to use AI to build your monitoring strategy:

- Data drift detection: Input data distribution changes (e.g., a new type of merchant becomes common). Prompt AI to write code that compares current data statistics to training data statistics daily.
- Concept drift: The relationship between features and fraud changes (fraudsters adapt). Monitor model accuracy on a rolling window. Ask AI to write an alerting system that triggers when accuracy drops below threshold.
- Latency monitoring: Ensure your model API responds within SLA (e.g., under 200ms for real-time fraud). Use AI to write a Prometheus metrics exporter for your FastAPI service.
- Volume monitoring: Alert if transaction volume drops 30% below normal — could indicate a data pipeline failure. AI can write Cloudwatch or Azure Monitor alert rules for this.

Monitoring Dashboard Prompt

```
"Design a monitoring dashboard for a real-time fraud detection AI system. List:
all metrics to track (with descriptions and alert thresholds), the ideal refresh
frequency for each metric, what actions to take when each alert fires, and
suggested visualisation types for each metric. Consider: model accuracy, data
quality, system performance, and business impact metrics."
```

6.6 Communicating Data Insights

Numbers without narrative are useless to business stakeholders. AI can help you write the story that makes your data analysis actionable:

```
"Based on these analysis results: [paste your metrics/findings], write a
3-paragraph executive summary that: 1) States the key finding in plain English,
2) Explains the business impact (use dollar amounts where possible), 3)
Recommends a specific action with expected outcome. Audience: VP of Risk
Management (financial background, not technical)."
```

■ **Fiserv Fast Start — Chapter 6**	Take any dataset you are currently working with — even a simple spreadsheet. Ask AI: 'What are the 5 most valuable analyses I could run on a dataset with these columns: [list your columns]? For each analysis, describe: what question it answers, what to look for, and the Python code to run it.' Then run at least one and share the result with your team.

CHAPTER 7

AI for App Building

Ship AI-Powered Features — From Prototype to Production in the Fiserv Environment

7.1 Building AI Applications

This is where everything comes together. As an AI Implementation Engineer, your primary output is working AI systems — APIs, services, pipelines, and features that run in production. AI tools transform how you build these systems: faster prototyping, better code quality, fewer bugs, and more time to focus on the hard architectural problems.

7.2 AI-Assisted Development Workflow

1. Understand the requirement	Use AI to clarify vague requirements: 'Rewrite this user story with clear acceptance criteria and list what technical components are needed.'
2. Design the architecture	Brainstorm with AI: 'Given these requirements, design a system architecture. What are the main components and how do they communicate?'
3. Write the code	AI generates boilerplate, complex algorithms, integration code. You add business logic and review everything.
4. Write tests	AI generates unit tests, integration tests, and edge cases you might miss.
5. Debug and review	Paste broken code: 'This code produces this error. Find the bug and fix it.'
6. Write documentation	AI documents your code, generates API docs, and writes the README.
7. Code review	AI acts as a first-pass reviewer: 'Review this code for bugs, security issues, and performance problems.'

7.3 Building a RAG System for Fiserv

The most common AI implementation you will build at Fiserv is a RAG (Retrieval-Augmented Generation) system — an AI that can answer questions based on Fiserv's internal documents. Here is the full implementation guide, step by step:

Step 1: Document Ingestion Pipeline

```
"Write a Python script that: 1) Reads all PDF files from an Azure Blob Storage
container, 2) Splits each PDF into chunks of 500 tokens with 50-token overlap,
3) Generates embeddings for each chunk using Azure OpenAI
text-embedding-ada-002, 4) Stores embeddings and metadata (source file, page
number, chunk text) in Pinecone. Include error handling, progress logging, and
cost estimation."
```

Step 2: Query and Retrieval

```
"Write a Python function that: 1) Takes a user question as input, 2) Generates
an embedding for the question using Azure OpenAI, 3) Queries Pinecone to find
the top 5 most similar document chunks, 4) Returns the chunks with their source
metadata."
```

Step 3: Answer Generation

```
"Write a Python function that: 1) Takes the user question and retrieved document
chunks, 2) Constructs a prompt that instructs the model to answer based only on
the provided context, 3) Calls Azure OpenAI GPT-4 to generate the answer, 4)
Returns the answer with citations (which documents were used)."
```

Step 4: FastAPI Service

```
"Wrap the above functions into a FastAPI service with: POST /query endpoint
(takes question, returns answer + sources), GET /health endpoint, POST /ingest
endpoint (triggers document ingestion), Rate limiting (100 requests/minute),
Authentication via Azure AD token, Structured logging using structlog,
OpenTelemetry tracing."
```

7.4 Building an AI Agent for Fiserv Tasks

AI agents are AI systems that can take actions — not just answer questions. They can call APIs, query databases, send notifications, and complete multi-step workflows. Here is how to build a simple agent for Fiserv:

```
"Build a LangGraph agent that can: 1) Look up a customer's recent transactions
(mock this with a function that returns sample data), 2) Calculate their current
month spending by category, 3) Compare to their 3-month average, 4) Generate a
personalised spending insight report. Use Claude as the LLM, implement proper
error handling, and add a human-in-the-loop step before sending any
notification."
```

7.5 Code Review and Security

Financial systems have high security requirements. AI is an excellent first-pass security reviewer — but not a replacement for your security team's review.

```
"Review this Python code for security vulnerabilities relevant to a financial
services API. Check for: SQL injection risks, authentication weaknesses,
sensitive data exposure (logging PII/PAN data), insecure dependencies, input
validation issues, and OWASP Top 10 vulnerabilities. For each issue found,
explain the risk and the fix."
```

7.6 Testing AI Systems

Testing AI systems requires different approaches than testing traditional software, because AI output is non-deterministic. AI helps you design robust test strategies:

- Functional tests: 'Generate 50 test questions for our document Q&A; system covering: questions it should answer well, edge cases, questions outside its scope, and adversarial inputs (attempts to make it give wrong answers).'
- Regression tests: 'Create a test harness that runs our RAG system against a golden set of 20 question-answer pairs and alerts if answer quality drops below a threshold.'
- Load tests: 'Write a Locust load test script for our fraud detection API. Simulate 500 concurrent users sending 1,000 transactions per second for 5 minutes.'
- Chaos tests: 'List 10 failure scenarios to test for our AI pipeline and the expected system behaviour in each scenario. Include: vector database unavailability, LLM API rate limiting, malformed input data, and upstream service outages.'

7.7 Deployment and Operations

Getting AI code into production at Fiserv requires containerisation, CI/CD, and careful monitoring. AI helps you build all of these:

```
"Write a production-ready Dockerfile for a FastAPI Python application that
serves an AI model. Include: multi-stage build for smaller image size, non-root
user for security, health check, proper signal handling, and environment
variable configuration. Also write a GitHub Actions CI/CD pipeline that: runs
tests, builds the Docker image, pushes to Azure Container Registry, and deploys
to Azure Container Apps."
```

■ **Fiserv Fast Start — Chapter 7**	Start your first AI application today. Use AI to build a simple RAG system that can answer questions about a document you care about (e.g., Fiserv's API documentation, an architecture document, or your team's runbook). You will need: Python, an OpenAI API key (use a personal one for testing), and ChromaDB (local vector database — no setup required). Prompt: 'Give me a minimal working Python script for a local RAG system using OpenAI embeddings and ChromaDB that I can run in under 10 minutes.' Then extend it.

CONCLUSION

Your 30-Day AI Action Plan at Fiserv

You have now covered all seven AI skill areas. The worst thing you can do is read this guide and do nothing. The second-worst thing is trying to do everything at once. Here is a structured 30-day plan to make these skills real in your daily work at Fiserv.

Week 1: Foundation (Days 1–7)

- Day 1: Set up access to your AI tool (Claude, Copilot, or ChatGPT). Write your first 5 prompts based on Chapter 1.
- Day 2: Use AI to brainstorm improvements to one real project you are working on.
- Day 3: Research one new technology your team is considering. Use the Orient → Deep Dive → Connect pattern from Chapter 3.
- Day 4: Rewrite a document you wrote last week using AI assistance. Compare the before and after.
- Day 5: Create a one-page training cheat sheet for your team using AI.
- Day 6–7: Run EDA on any dataset you have access to using AI-generated Python code.

Week 2: Apply (Days 8–14)

- Day 8: Write your first AI-assisted Technical Design Document for a real project.
- Day 9: Ask AI to generate test cases for a feature you are currently building.
- Day 10: Build a minimal RAG system following Chapter 7. Even locally, even for a tiny document.
- Day 11: Use AI to write a business case for one AI initiative you believe Fiserv should pursue.
- Day 12: Create a monitoring plan for one AI system you own or work on.
- Day 13: Have AI review your code for security issues. Fix at least two of them.
- Day 14: Share something you learned with your team (a prompt, a technique, a tool).

Week 3: Deepen (Days 15–21)

- Day 15: Spend 2 hours building something more ambitious — a simple AI agent or a more complex RAG pipeline.
- Day 16: Use AI to prepare for an upcoming meeting (brainstorm questions, generate an agenda).
- Day 17: Research a compliance/regulatory topic relevant to an AI system you work on.
- Day 18: Create training materials for end users of an AI system at Fiserv.
- Day 19: Build a data pipeline analysis script using AI-generated code.
- Day 20: Conduct an AI-assisted incident post-mortem on a recent issue (even a minor one).
- Day 21: Reflect: which AI skill has had the biggest impact on your work so far?

Week 4: Scale (Days 22–30)

- Day 22–24: Start a meaningful AI project — something that could save your team real time or improve quality.
- Day 25: Present what you have built to your team. Get feedback.
- Day 26: Document your learnings. Create a guide for your team based on what worked.
- Day 27–28: Help one colleague get started with AI. Teaching accelerates your own learning.
- Day 29: Set your 90-day AI goals. What AI project do you want to complete? What skill do you want to master?
- Day 30: Celebrate. You have gone from knowing where to start to actually doing it.

The AI Engineer's Mindset

The engineers who will thrive in the AI era are not necessarily the ones who know the most about AI. They are the ones who are curious, who experiment, who share what they learn, and who stay humble enough to keep learning.

- Be curious: Every time you do something repetitive, ask 'Could AI help with this?'
- Be experimental: Not every AI experiment will work. That is fine. Fail fast, learn, move on.
- Be honest: Always tell stakeholders when AI was used and what its limitations are.
- Be ethical: Think about the impact of AI systems on the customers Fiserv serves. They trust us.
- Be collaborative: Share your AI techniques with your team. Rising tide lifts all boats.

The real AI divide is not between people who are smart and people who are not.
It is between people who started learning AI today and people who will start tomorrow.

You have already started. Keep going.